

Datasets and cross-centre evaluation protocol

Dataset	Images	Role in this study	Protocol
Kvasir-SEG	1,000	Train / validation / test	Seen
CVC-ClinicDB	612	Train / validation / test	Seen
CVC-ColonDB	380	Test only	Unseen
ETIS-LaribPolypDB	196	Test only	Unseen
CVC-300 (EndoScene)	60	Test only	Unseen

Training pool = Kvasir-SEG + CVC-ClinicDB (1,612 images), split 80/20 into 1,289 training / 323 validation, at 352×352 px.
The three unseen datasets are held out entirely for cross-centre generalization testing (never seen during training).